Name:Falguni Gaikwad

Rollno:J077

GridSearchCV API

GridSearchCV is a library function that is a member of sklearn's model_selection package. It helps to loop through predefined hyperparameters and fit your estimator (model) on your training set. So, in the end, you can select the best parameters from the listed hyperparameters.

API:

Exhaustive search over specified parameter values for an estimator.

Important members are fit, predict.

GridSearchCV implements a "fit" and a "score" method. It also implements "score_samples", "predict", "predict_proba", "decision_function", "transform" and "inverse_transform" if they are implemented in the estimator used.

The parameters of the estimator used to apply these methods are optimized by cross-validated grid-search over a parameter grid.

CODE:

class sklearn.model_selection.GridSearchCV(estimator, param_grid, *, scoring=None, n_jobs=None, refit=True, cv=None, verbose=0, pre_dispatch='2*n_jobs', error_score=nan, return_train_score=False)

PARAMETERS:

estimator: estimator object.

This is assumed to implement the scikit-learn estimator interface. Either estimator needs to provide a score function, or scoring must be passed.

param_grid: dict or list of dictionaries

Dictionary with parameters names (str) as keys and lists of parameter settings to try as values, or a list of such dictionaries, in which case the grids spanned by each dictionary in the list are explored.

scoring: str, callable, list, tuple or dict, default=None

Strategy to evaluate the performance of the cross-validated model on the test set.

1. If scoring represents a single score, one can use:
    a. a single string;

b. a callable that returns a single value.

2. If scoring represents multiple scores, one can use:
    a. a list or tuple of unique strings;
    b. a callable returning a dictionary where the keys are the metric names and the values are the metric scores;
    c. a dictionary with metric names as keys and callables a values.

n_jobs: int, default=None

Number of jobs to run in parallel.

refit: bool, str, or callable, default=True

Refit an estimator using the best found parameters on the whole dataset.

cv: int, cross-validation generator or an iterable, default=None

Determines the cross-validation splitting strategy.

verbose: int

Controls the verbosity: the higher, the more messages.

>1 : the computation time for each fold and parameter candidate is displayed;

>2 : the score is also displayed;

>3 : the fold and candidate parameter indexes are also displayed together with the starting time of the computation.

pre_dispatch: int, or str, default=n_jobs

Controls the number of jobs that get dispatched during parallel execution.

error_score'raise' or numeric, default=np.nan

Value to assign to the score if an error occurs in estimator fitting. If set to 'raise', the error is raised.

return_train_scorebool, default=False

If False, the cv_results_ attribute will not include training scores.

ATTRIBUTES:

cv_results_:dict of numpy (masked) ndarrays

A dict with keys as column headers and values as columns, that can be imported into a pandas DataFrame.

best_estimator_: estimator

Estimator that was chosen by the search, i.e. estimator which gave highest score (or smallest loss if specified) on the left out data. Not available if refit=False.

best_score_: float

Mean cross-validated score of the best_estimator.

best_params_: dict

Parameter setting that gave the best results on the hold out data.

best_index_: int

The index (of the cv_results_ arrays) which corresponds to the best candidate parameter setting..

scorer_: function or a dict

Scorer function used on the held out data to choose the best parameters for the model.

n_splits_: int

The number of cross-validation splits (folds/iterations).

refit_time_: float

Seconds used for refitting the best model on the whole dataset.

multimetric_: bool

Whether or not the scorers compute several metrics.