

Sci-kit learn API

1. Linear Regression:

- **Code:** `sklearn.linear_model.LinearRegression(*, fit_intercept=True, normalize=False, copy_X=True, n_jobs=None, positive=False)`

`fit_intercept`bool, *default=True*. Whether to calculate the intercept for this model.

`normalize`bool, *default=False*. This parameter is ignored when `fit_intercept` is set to False.

`copy_X`bool, *default=True*. If True, X will be copied; else, it may be overwritten.

`n_jobs`int, *default=None*. The number of jobs to use for the computation.

`positive`bool, *default=False*. When set to True, forces the coefficients to be positive.

- **“LinearRegression”** fits a linear model with coefficients $w=(w_1, \dots, w_p)$ to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation. . In its fit method arrays X, y and will store the coefficients of the linear model in its `coef_` member.

2. Logistic Regression:

- **Code:** `sklearn.linear_model.LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)`

`penalty`{‘l1’, ‘l2’, ‘elasticnet’, ‘none’}, *default=‘l2’*

Used to specify the norm used in the penalization.

`dual`bool, *default=False*

Dual or primal formulation.

`tol`float, *default=1e-4*

Tolerance for stopping criteria.

`C`float, *default=1.0*

Inverse of regularization strength; must be a positive float.

`fit_intercept`bool, *default=True*

Specifies if a constant (a.k.a. bias or intercept) should be added to the decision function.

`intercept_scaling`float, *default=1*

Useful only when the solver 'liblinear' is used and `self.fit_intercept` is set to True.

`class_weight`dict or 'balanced', *default=None*

Weights associated with classes in the form {`class_label`: `weight`}. If not given, all classes are supposed to have weight one.

`random_state`int, *RandomState instance, default=None*

Used when `solver == 'sag', 'saga' or 'liblinear'` to shuffle the data.

`solver`='lbfgs'

'lbfgs' handle multinomial loss

`max_iter`int, *default=100*

Maximum number of iterations taken for the solvers to converge.

`multi_class`{'auto', 'ovr', 'multinomial'}, *default='auto'*

If the option chosen is 'ovr', then a binary problem is fit for each label.

`verbose`int, *default=0*

For the liblinear and lbfgs solvers set verbose to any positive number for verbosity.

`warm_start`bool, *default=False*

When set to True, reuse the solution of the previous call to fit as initialization

`n_jobs`int, *default=None*

Number of CPU cores used when parallelizing over classes if `multi_class='ovr'`.

`l1_ratio`float, *default=None*

The Elastic-Net mixing parameter, with $0 \leq l1_ratio \leq 1$.

- **“LogisticRegression”**, despite its name, is a linear model for classification rather than regression. Logistic regression is also known in the literature as logit regression, maximum-entropy classification (MaxEnt) or the log-linear classifier. In this model, the probabilities describing the possible outcomes of a single trial are modeled using a logistic function.

3. Ridge:

Code: `sklearn.linear_model.Ridge(alpha=1.0, *, fit_intercept=True, normalize=False, copy_X=True, max_iter=None, tol=0.001, solver='auto', random_state=None)`

`alpha`{float, ndarray of shape (n_targets,)}, *default=1.0*

Regularization strength; must be a positive float.

`fit_intercept`bool, *default=True*

Whether to fit the intercept for this model.

`normalize`bool, *default=False*

This parameter is ignored when `fit_intercept` is set to `False`.

`copy_Xbool, default=True`

If `True`, `X` will be copied; else, it may be overwritten.

`max_iterint, default=None`

Maximum number of iterations for conjugate gradient solver.

`tolfloat, default=1e-3`

Precision of the solution.

`solver{'auto'}`

'auto' chooses the solver automatically based on the type of data.

- **"Ridge"** regression addresses some of the problems of '*Ordinary Least Squares*' by imposing a penalty on the size of the coefficients. The ridge coefficients minimize a penalized residual sum of squares. The complexity parameter $\alpha > 0$ controls the amount of shrinkage: the larger the value of α , the greater the amount of shrinkage and thus the coefficients become more robust to collinearity. In its fit method arrays `X`, `y` and will store the coefficients of the linear model in its `coef_` member.

4. Lasso:

- **Code:** `sklearn.linear_model.Lasso(alpha=1.0, *, fit_intercept=True, normalize=False, precompute=False, copy_X=True, max_iter=1000, tol=0.0001, warm_start=False, positive=False, random_state=None, selection='cyclic')`

`alphafloat, default=1.0`

Constant that multiplies the L1 term. Defaults to 1.0. $\alpha = 0$ is equivalent to an ordinary least square, solved by the `LinearRegression` object.

`fit_interceptbool, default=True`

Whether to calculate the intercept for this model.

`normalizebool, default=False`

This parameter is ignored when `fit_intercept` is set to `False`.

`precomputebool or array-like of shape (n_features, n_features), default=False`

Whether to use a precomputed Gram matrix to speed up calculations.

`copy_Xbool, default=True`

If `True`, `X` will be copied; else, it may be overwritten.

`max_iterint, default=1000`

The maximum number of iterations.

`tolfloat, default=1e-4`

The tolerance for the optimization: if the updates are smaller than `tol`, the optimization code checks the dual gap for optimality and continues until it is smaller than `tol`.

`warm_startbool, default=False`

When set to True, reuse the solution of the previous call to fit as initialization, otherwise, just erase the previous solution.

positivebool, default=False

When set to True, forces the coefficients to be positive.

random_stateint, RandomState instance, default=None

The seed of the pseudo random number generator that selects a random feature to update.

selection{'cyclic', 'random'}, default='cyclic'

If set to 'random', a random coefficient is updated every iteration rather than looping over features sequentially by default.

- The “**Lasso**” is a linear model that estimates sparse coefficients. It is useful in some contexts due to its tendency to prefer solutions with fewer non-zero coefficients, effectively reducing the number of features upon which the given solution is dependent.